

MACHINE LEARNING MAJOR PROJECT

Problem Statement :-

Take any Dataset of your choice, perform EDA (Exploratory Data Analysis) and apply a suitable Classifier, Regressor or Clusterer and calculate the accuracy of the model.

I have chosen the 'Salary data' dataset and I will be performing Linear regression on it.

Link of the dataset - https://github.com/ameenmanna8824/DATASETS/blob/main/Salary_Data.csv

Raw - https://raw.githubusercontent.com/ameenmanna8824/DATASETS/main/Salary_Data.csv

Creating dataframe

```
In [53]: import pandas as pd
df=pd.read_csv('https://raw.githubusercontent.com/ameenmanna8824/DATASETS/main/Salary_Data.csv')
df
```

Out[53]:

	YearsExperience	Salary
0	1.1	39343.0
1	1.3	46205.0
2	1.5	37731.0
3	2.0	43525.0
4	2.2	39891.0
5	2.9	56642.0
6	3.0	60150.0
7	3.2	54445.0
8	3.2	64445.0
9	3.7	57189.0
10	3.9	63218.0
11	4.0	55794.0
12	4.0	56957.0
13	4.1	57081.0
14	4.5	61111.0
15	4.9	67938.0
16	5.1	66029.0
17	5.3	83088.0
18	5.9	81363.0
19	6.0	93940.0
20	6.8	91738.0
21	7.1	98273.0
22	7.9	101302.0
23	8.2	113812.0

	YearsExperience	Salary
24	8.7	109431.0
25	9.0	105582.0
26	9.5	116969.0
27	9.6	112635.0
28	10.3	122391.0
29	10.5	121872.0

Exploratory Data Analysis (EDA)

```
In [54]: df.shape
```

```
Out[54]: (30, 2)
```

```
In [55]: df.YearsExperience.nunique()
```

```
Out[55]: 28
```

```
In [56]: df.Salary.nunique()
```

```
Out[56]: 30
```

```
In [57]: df.Salary.unique()
```

```
Out[57]: array([ 39343., 46205., 37731., 43525., 39891., 56642., 60150.,
        54445., 64445., 57189., 63218., 55794., 56957., 57081.,
        61111., 67938., 66029., 83088., 81363., 93940., 91738.,
        98273., 101302., 113812., 109431., 105582., 116969., 112635.,
        122391., 121872.])
```

```
In [58]: df.YearsExperience.unique()
```

```
Out[58]: array([ 1.1,  1.3,  1.5,  2. ,  2.2,  2.9,  3. ,  3.2,  3.7,  3.9,  4. ,
          4.1,  4.5,  4.9,  5.1,  5.3,  5.9,  6. ,  6.8,  7.1,  7.9,  8.2,
          8.7,  9. ,  9.5,  9.6, 10.3, 10.5])
```

Salary analysis

```
In [59]: # Salary above 40k
```

```
df['Salary above 40k'] = df['Salary'] > 40000
df.groupby('Salary above 40k').size()
```

```
Out[59]: Salary above 40k
False      3
True       27
dtype: int64
```

```
In [60]: # Salary above 70k
```

```
df['Salary above 70k'] = df['Salary'] > 70000
df.groupby('Salary above 70k').size()
```

```
Out[60]: Salary above 70k
False     17
True      13
dtype: int64
```

```
In [62]: # Salary above 100k
```

```
df['Salary above 100k'] = df['Salary'] > 100000
df.groupby('Salary above 100k').size()
```

```
Out[62]: Salary above 100k
False     22
True       8
dtype: int64
```

From the above cell blocks, we can deduce that

- Employees with salary above 40,000 : 27 / 30

- Employees with salary above 70,000 : 13 / 30
 - Employees with salary above 100,000 : 8 / 30
-

Years of exp. analysis

```
In [63]: # Years of Experience above 2

df['Year of exp. > 2'] = df['YearsExperience'] > 2
df.groupby('Year of exp. > 2').size()
```

```
Out[63]: Year of exp. > 2
False      4
True      26
dtype: int64
```

```
In [64]: # Years of Experience above 5

df['Year of exp. > 5'] = df['YearsExperience'] > 5
df.groupby('Year of exp. > 5').size()
```

```
Out[64]: Year of exp. > 5
False     16
True      14
dtype: int64
```

```
In [65]: # Years of Experience above 8

df['Year of exp. > 8'] = df['YearsExperience'] > 8
df.groupby('Year of exp. > 8').size()
```

```
Out[65]: Year of exp. > 8
False     23
True       7
dtype: int64
```

```
In [66]: # Years of Experience above 10

df['Year of exp. > 10'] = df['YearsExperience'] > 10
df.groupby('Year of exp. > 10').size()
```

```
Out[66]: Year of exp. > 10
False    28
True      2
dtype: int64
```

From the above cell blocks, we can deduce that

- Employees with years of exp. above 2 : 26 / 30
 - Employees with years of exp. above 5 : 14 / 30
 - Employees with years of exp. above 8 : 7 / 30
 - Employees with years of exp. above 10 : 2 / 30
-

Min. and Max. years of experience and salary

```
In [67]: df.YearsExperience.min()
```

```
Out[67]: 1.1
```

```
In [68]: df.YearsExperience.max()
```

```
Out[68]: 10.5
```

```
In [69]: df.Salary.min()
```

```
Out[69]: 37731.0
```

```
In [70]: df.Salary.max()
```

```
Out[70]: 122391.0
```

From the above cells, we can deduce that

- Min. Years of exp. : 1.1 yrs
 - Max. Years of exp. : 10.5 yrs
 - Min. Salary : 37731.0
 - Max. Salary : 122391.0
-
-

Data Visualisation

```
In [71]: import matplotlib.pyplot as plt
plt.scatter(df['YearsExperience'], df['Salary'])
plt.title("YEARS OF EXPERIENCE VS SALARY")
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
```

```
Out[71]: Text(0, 0.5, 'Salary')
```



```
Out[72]:
```

0	1.1
1	1.3
2	1.5
3	2.0
4	2.2
5	2.9
6	3.0
7	3.2
8	3.2
9	3.7
10	3.9
11	4.0
12	4.0
13	4.1
14	4.5
15	4.9
16	5.1
17	5.3
18	5.9
19	6.0
20	6.8
21	7.1
22	7.9
23	8.2
24	8.7
25	9.0
26	9.5
27	9.6
28	10.3
29	10.5

Name: YearsExperience, dtype: float64

```
In [73]: x = df.iloc[:29,:1].values
x
```

```
Out[73]: array([[ 1.1],
 [ 1.3],
 [ 1.5],
 [ 2. ],
 [ 2.2],
 [ 2.9],
 [ 3. ],
 [ 3.2],
 [ 3.2],
 [ 3.7],
 [ 3.9],
 [ 4. ],
 [ 4. ],
 [ 4.1],
 [ 4.5],
 [ 4.9],
 [ 5.1],
 [ 5.3],
 [ 5.9],
 [ 6. ],
 [ 6.8],
 [ 7.1],
 [ 7.9],
 [ 8.2],
 [ 8.7],
 [ 9. ],
 [ 9.5],
 [ 9.6],
 [10.3]])
```

```
In [74]: y = df.iloc[:,1]
y
```

```
Out[74]: 0      39343.0
          1      46205.0
          2      37731.0
          3      43525.0
          4      39891.0
          5      56642.0
          6      60150.0
          7      54445.0
          8      64445.0
          9      57189.0
         10      63218.0
         11      55794.0
         12      56957.0
         13      57081.0
         14      61111.0
         15      67938.0
         16      66029.0
         17      83088.0
         18      81363.0
         19      93940.0
         20      91738.0
         21      98273.0
         22     101302.0
         23     113812.0
         24     109431.0
         25     105582.0
         26     116969.0
         27     112635.0
         28     122391.0
         29     121872.0
Name: Salary, dtype: float64
```

```
In [75]: y = df.iloc[1:30,1].values
          y
```

```
Out[75]: array([ 46205.,  37731.,  43525.,  39891.,  56642.,  60150.,  54445.,
                64445.,  57189.,  63218.,  55794.,  56957.,  57081.,  61111.,
                67938.,  66029.,  83088.,  81363.,  93940.,  91738.,  98273.,
                101302., 113812., 109431., 105582., 116969., 112635., 122391.,
                121872.])
```

Running Regressor

```
In [76]: from sklearn.linear_model import LinearRegression
model = LinearRegression()
```

Fitting the model

```
In [77]: model.fit(x,y)
```

```
Out[77]: LinearRegression()
```

Predicting the output

```
In [78]: # Predicted Output Values
```

```
y_pred = model.predict(x)
y_pred
```

```
Out[78]: array([ 38160.81764789,  40099.42155936,  42038.02547083,  46884.53524951,
  48823.13916098,  55608.25285114,  56577.55480687,  58516.15871835,
  58516.15871835,  63362.66849703,  65301.2724085 ,  66270.57436423,
  66270.57436423,  67239.87631997,  71117.08414292,  74994.29196586,
  76932.89587733,  78871.4997888 ,  84687.31152322,  85656.61347896,
  93411.02912485,  96318.93499206, 104073.35063794, 106981.25650515,
 111827.76628383, 114735.67215104, 119582.18192972, 120551.48388546,
 127336.59757561])
```

```
In [79]: # Actual Output Values
```

```
y
```

```
Out[79]: array([ 46205.,  37731.,  43525.,  39891.,  56642.,  60150.,  54445.,
  64445.,  57189.,  63218.,  55794.,  56957.,  57081.,  61111.,
  67938.,  66029.,  83088.,  81363.,  93940.,  91738.,  98273.,
 101302., 113812., 109431., 105582., 116969., 112635., 122391.,
 121872.])
```

Conclulsion

When we compare y and pred_y , we come to know that there is a huge difference. This huge difference does not mean that my model has predicted wrong. It only means that my model is NOT LINEAR / LESS LINEAR.

Individual Prediction

```
In [80]: model.predict([[5]])
```

```
Out[80]: array([75963.5939216])
```

Cross-verifiacion

$$y = mx + c$$

```
In [81]: m = model.coef_  
m
```

```
Out[81]: array([9693.01955736])
```

```
In [82]: c = model.intercept_  
c
```

```
Out[82]: 27498.496134789544
```

```
In [83]: m * 5 + c
```

```
Out[83]: array([75963.5939216])
```

Accuracy of the model (R-squared score)

Since the answer to the above cell block is the same as the 'Individual Prediction' block, our model has predicted correct.

Furthermore, i will find the R-squared score of the model.

(The closer the R-squared score is to 1, the more accurate the model.)

```
In [84]: from sklearn.metrics import r2_score
```

```
r2_score = r2_score(y, y_pred)
print("R-squared score:", r2_score)
```

```
R-squared score: 0.9472760886559446
```

Since the R-squared score is very close to 1, we can say that our model is fairly accurate.

Final Visualisation

```
In [85]: plt.scatter(x,y,c='blue') # Actual output values
plt.plot(x,y_pred,color='red') # Predicted output values

# the scattered points will be the actual values and the line will represent the predicted values

plt.title("BEST FIT LINE")
plt.xlabel("Years of experience")
plt.ylabel("Salary")
```

```
Out[85]: Text(0, 0.5, 'Salary')
```

